

Tutorial 6 - Deep Learning

Kevin Dick, PhDc Biomedical Engineering
Carleton University

Friday 16th October, 2020

Asynchronous Tutorial


To watch and following along at your leisure

ML Weekly

Recent news events from the ML community

1. (ML) Google Colab: How does it compare to a GPU-enabled laptop?

Fashion-MNIST - Training Time (seconds) Comparison



Sum of Seconds for each Machine.

ML Weekly

1. **(ML)** Google Colab: How does it compare to a GPU-enabled laptop?
2. **(Vision)** "NeurIPS 2020 Workshop — Indie GAN Interpolation Method Turns Selfies Into Cartoon Characters



ML Weekly

1. **(ML)** Google Colab: How does it compare to a GPU-enabled laptop?
2. **(Vision)** "NeurIPS 2020 Workshop — Indie GAN Interpolation Method Turns Selfies Into Cartoon Characters




ML Weekly

1. **(ML)** Google Colab: How does it compare to a GPU-enabled laptop?
2. **(Vision)** "NeurIPS 2020 Workshop — Indie GAN Interpolation Method Turns Selfies Into Cartoon Characters
3. **(ML)** A radical new technique lets AI learn with practically no data



ML Weekly

1. **(ML)** Google Colab: How does it compare to a GPU-enabled laptop?
2. **(Vision)** "NeurIPS 2020 Workshop — Indie GAN Interpolation Method Turns Selfies Into Cartoon Characters
3. **(ML)** A radical new technique lets AI learn with practically no data
4. **(Vision)** Benchmark analysis of representative deep neural network architectures



Tutorial Intuition

Building an Intuition for the Concepts of this Tutorial

Introducing the Keras API: Deep Network in < 30 Lines

The core data structures of Keras are **layers** and **models**. The simplest type of model is the **Sequential model**, a linear stack of layers.

Algorithm 1 Pseudocode for Building a Deep Learning (Vision) Model

Input: model parameters

Output: model object implementing architecture

- 1: initialized model object, m
 - 2: $m \leftarrow$ add **input** layer
 - 3: **for each** hidden layer, h_i , in $\{h_1, h_2, \dots, h_n\}$ **do**
 - 4: $m \leftarrow$ add a **convolutional** layer
 - 5: $m \leftarrow$ add an **activation** layer (optional)
 - 6: $m \leftarrow$ add a **normalization** layer (optional)
 - 7: $m \leftarrow$ add a **pooling** layer (optional)
 - 8: $m \leftarrow$ add a **dropout** layer (optional)
 - 9: **end for**
 - 10: $m \leftarrow$ add **output** layer
 - 11: return m
-

Introducing the Keras API

```
1 from tensorflow.keras.models import Sequential
2 from tensorflow.keras.layers import Dense
3
4 model = Sequential() # a linear stack of layers
5
6 # stacking layers is as easy as .add():
7 #
8 #
9
10 # configure model learning process with .compile():
11 #
12 #
13 #
14
15 # (Optional) further configure your optimizer.
16 #
17 #
18 #
19
20 # x_train and y_train are Numpy arrays akin to Scikit-Learn API.
21 #
22
23 # Evaluate your test loss and metrics in one line:
24 #
25
26 # Or generate predictions on new data:
27 #
```

Introducing the Keras API

```
1 from tensorflow.keras.models import Sequential
2 from tensorflow.keras.layers import Dense
3
4 model = Sequential() # a linear stack of layers
5
6 # stacking layers is as easy as .add():
7 model.add(Dense(units=64, activation='relu'))
8 model.add(Dense(units=10, activation='softmax'))
9
10 # configure model learning process with .compile():
11 #
12 #
13 #
14
15 # (Optional) further configure your optimizer.
16 #
17 #
18 #
19
20 # x_train and y_train are Numpy arrays akin to Scikit-Learn API.
21 #
22
23 # Evaluate your test loss and metrics in one line:
24 #
25
26 # Or generate predictions on new data:
27 #
```

Introducing the Keras API

```
1 from tensorflow.keras.models import Sequential
2 from tensorflow.keras.layers import Dense
3
4 model = Sequential() # a linear stack of layers
5
6 # stacking layers is as easy as .add():
7 model.add(Dense(units=64, activation='relu'))
8 model.add(Dense(units=10, activation='softmax'))
9
10 # configure model learning process with .compile():
11 model.compile(loss='categorical_crossentropy',
12               optimizer='sgd',
13               metrics=['accuracy'])
14
15 # (Optional) further configure your optimizer.
16 model.compile(loss=keras.losses.categorical_crossentropy,
17               optimizer=keras.optimizers.SGD(learning_rate=0.01,
18                                                momentum=0.9, nesterov=True))
19
20 # x_train and y_train are Numpy arrays akin to Scikit-Learn API.
21 #
22
23 # Evaluate your test loss and metrics in one line:
24 #
25
26 # Or generate predictions on new data:
27 #
```

Introducing the Keras API

```
1 from tensorflow.keras.models import Sequential
2 from tensorflow.keras.layers import Dense
3
4 model = Sequential() # a linear stack of layers
5
6 # stacking layers is as easy as .add():
7 model.add(Dense(units=64, activation='relu'))
8 model.add(Dense(units=10, activation='softmax'))
9
10 # configure model learning process with .compile():
11 model.compile(loss='categorical_crossentropy',
12               optimizer='sgd',
13               metrics=['accuracy'])
14
15 # (Optional) further configure your optimizer.
16 model.compile(loss=keras.losses.categorical_crossentropy,
17               optimizer=keras.optimizers.SGD(learning_rate=0.01,
18                                                momentum=0.9, nesterov=True))
19
20 # x_train and y_train are Numpy arrays akin to Scikit-Learn API.
21 model.fit(x_train, y_train, epochs=5, batch_size=32)
22
23 # Evaluate your test loss and metrics in one line:
24 #
25
26 # Or generate predictions on new data:
27 #
```

Introducing the Keras API

```
1 from tensorflow.keras.models import Sequential
2 from tensorflow.keras.layers import Dense
3
4 model = Sequential() # a linear stack of layers
5
6 # stacking layers is as easy as .add():
7 model.add(Dense(units=64, activation='relu'))
8 model.add(Dense(units=10, activation='softmax'))
9
10 # configure model learning process with .compile():
11 model.compile(loss='categorical_crossentropy',
12               optimizer='sgd',
13               metrics=['accuracy'])
14
15 # (Optional) further configure your optimizer.
16 model.compile(loss=keras.losses.categorical_crossentropy,
17               optimizer=keras.optimizers.SGD(learning_rate=0.01,
18                                                 momentum=0.9, nesterov=True))
19
20 # x_train and y_train are Numpy arrays akin to Scikit-Learn API.
21 model.fit(x_train, y_train, epochs=5, batch_size=32)
22
23 # Evaluate your test loss and metrics in one line:
24 loss_and_metrics = model.evaluate(x_test, y_test, batch_size=128)
25
26 # Or generate predictions on new data:
27 classes = model.predict(x_test, batch_size=128)
```

Into the Notebooks we Go...

We will cover one new notebook today!

1. Tutorial 6 - Deep Learning

Tutorial 6 - Deep Learning

Kevin Dick, PhDc Biomedical Engineering
Carleton University

Friday 16th October, 2020